


```
In [1]: 1 # set tf 1.x for colab
        2 %tensorflow_version 1.x
```

UsageError: Line magic function `%tensorflow_version` not found.

```
In [2]: 1 import warnings
        2 warnings.filterwarnings('ignore', category=DeprecationWarning)
        3 warnings.filterwarnings('ignore', category=FutureWarning)
```

MNIST digits classification with TensorFlow



```
In [3]: 1 import numpy as np
        2 from sklearn.metrics import accuracy_score
        3 from matplotlib import pyplot as plt
        4 %matplotlib inline
        5 import tensorflow as tf
        6 print("We're using TF", tf.__version__)
        7
        8 import sys
        9 sys.path.append("../..")
       10 import grading
       11
       12 import matplotlib_utils
       13 from importlib import reload
       14 reload(matplotlib_utils)
       15
       16 import grading_utils
       17 reload(grading_utils)
       18
       19 import keras_utils
       20 from keras_utils import reset_tf_session
```

We're using TF 1.14.0

Using TensorFlow backend.

Fill in your Coursera token and email

To successfully submit your answers to our grader, please fill in your Coursera submission token and email

```
In [3]: 1 grader = grading.Grader(assignment_key="XtD7ho3TEeiHQBLejjYAA",
        2                     all_parts=["9XaAS", "vmogZ", "RMv95", "i8bgs", "rE763"])
```

```
In [4]: 1 # token expires every 30 min
        2 COURSERA_TOKEN = "xsKaGouklo4Fsk8K"
        3 COURSERA_EMAIL = "lxwvictor@gmail.com"
```


Look at the data

In this task we have 50000 28x28 images of digits from 0 to 9. We will train a classifier on this data.


```
In [5]: 1 import preprocessed_mnist
        2 X_train, y_train, X_val, y_val, X_test, y_test = preprocessed_mnist.load_dataset()
```

```
In [6]: 1 # X contains rgb values divided by 255
2 print("X_train [shape %s] sample patch:\n" % (str(X_train.shape)), X_train[1, 15:20, 5:10])
3 print("A closeup of a sample patch:")
4 plt.imshow(X_train[1, 15:20, 5:10], cmap="Greys")
5 plt.show()
6 print("And the whole sample:")
7 plt.imshow(X_train[1], cmap="Greys")
8 plt.show()
9 print("y_train [shape %s] 10 samples:\n" % (str(y_train.shape)), y_train[:10])
```

```
X_train [shape (50000, 28, 28)] sample patch:
[[0.      0.29803922 0.96470588 0.98823529 0.43921569]
 [0.      0.33333333 0.98823529 0.90196078 0.09803922]
 [0.      0.33333333 0.98823529 0.8745098  0.      ]
 [0.      0.33333333 0.98823529 0.56862745 0.      ]
 [0.      0.3372549  0.99215686 0.88235294 0.      ]]
A closeup of a sample patch:
```



And the whole sample:



```
y_train [shape (50000,)] 10 samples:
[5 0 4 1 9 2 1 3 1 4]
```

Linear model

Your task is to train a linear classifier $\vec{x} \rightarrow y$ with SGD using TensorFlow.

You will need to calculate a logit (a linear transformation) z_k for each class:

$$z_k = \vec{x} \cdot \vec{w}_k + b_k \quad k = 0..9$$

And transform logits z_k to valid probabilities p_k with softmax:

$$p_k = \frac{e^{z_k}}{\sum_{i=0}^9 e^{z_i}} \quad k = 0..9$$

We will use a cross-entropy loss to train our multi-class classifier:

$$\text{cross-entropy}(y, p) = - \sum_{k=0}^9 \log(p_k)[y = k]$$

where

$$[x] = \begin{cases} 1, & \text{if } x \text{ is true} \\ 0, & \text{otherwise} \end{cases}$$

Cross-entropy minimization pushes p_k close to 1 when $y = k$, which is what we want.

Here's the plan:

- Flatten the images (28x28 -> 784) with `X_train.reshape((X_train.shape[0], -1))` to simplify our linear model implementation
- Use a matrix placeholder for flattened `X_train`
- Convert `y_train` to one-hot encoded vectors that are needed for cross-entropy
- Use a shared variable `W` for all weights (a column \vec{w}_k per class) and `b` for all biases.
- Aim for ~0.93 validation accuracy

```
In [7]: 1 X_train_flat = X_train.reshape((X_train.shape[0], -1))
2 print(X_train_flat.shape)
3
4 X_val_flat = X_val.reshape((X_val.shape[0], -1))
5 print(X_val_flat.shape)
```

```
(50000, 784)
(10000, 784)
```

```
In [8]: 1 import keras
2
3 y_train_oh = keras.utils.to_categorical(y_train, 10)
4 y_val_oh = keras.utils.to_categorical(y_val, 10)
5
6 print(y_train_oh.shape)
7 print(y_train_oh[:3], y_train[:3])
```

```
(50000, 10)
[[0. 0. 0. 0. 0. 1. 0. 0. 0. 0.]
 [1. 0. 0. 0. 0. 0. 0. 0. 0. 0.]
 [0. 0. 0. 0. 1. 0. 0. 0. 0. 0.]] [5 0 4]
```

```
In [9]: 1 # run this again if you remake your graph
2 s = reset_tf_session()
```

WARNING:tensorflow:From ../../keras_utils.py:68: The name tf.get_default_session is deprecated. Please use tf.compat.v1.get_default_session instead.

WARNING:tensorflow:From /Users/Victor/anaconda3/envs/tfspark/lib/python3.7/site-packages/keras/backend/tensorflow_backend.py:95: The name tf.reset_default_graph is deprecated. Please use tf.compat.v1.reset_default_graph instead.

WARNING:tensorflow:From /Users/Victor/anaconda3/envs/tfspark/lib/python3.7/site-packages/keras/backend/tensorflow_backend.py:98: The name tf.placeholder_with_default is deprecated. Please use tf.compat.v1.placeholder_with_default instead.

WARNING:tensorflow:From /Users/Victor/anaconda3/envs/tfspark/lib/python3.7/site-packages/keras/backend/tensorflow_backend.py:102: The name tf.get_default_graph is deprecated. Please use tf.compat.v1.get_default_graph instead.

WARNING:tensorflow:From ../../keras_utils.py:75: The name tf.ConfigProto is deprecated. Please use tf.compat.v1.ConfigProto instead.

```
In [10]: 1 # Model parameters: W and b
2 ### YOUR CODE HERE ### tf.get_variable(...) with shape[0] = 784
3 W = tf.get_variable('W', shape=(784, 10), dtype=tf.float32)
4 ### YOUR CODE HERE ### tf.get_variable(...)
5 b = tf.get_variable('b', shape=(10), dtype=tf.float32)
```

WARNING:tensorflow:From /Users/Victor/anaconda3/envs/tfspark/lib/python3.7/site-packages/tensorflow/python/ops/init_ops.py:1251: calling VarianceScaling.__init__ (from tensorflow.python.ops.init_ops) with dtype is deprecated and will be removed in a future version.

Instructions for updating:

Call initializer instance with the dtype argument instead of passing it to the constructor

```
In [11]: 1 # Placeholders for the input data
2 ### YOUR CODE HERE ### tf.placeholder(...) for flat X with shape[0] = None for any batch size
3 input_X = tf.placeholder(tf.float32, shape=(None, 784))
4 ### YOUR CODE HERE ### tf.placeholder(...) for one-hot encoded true labels
5 input_y = tf.placeholder(tf.int32, shape=(None, 10))
```

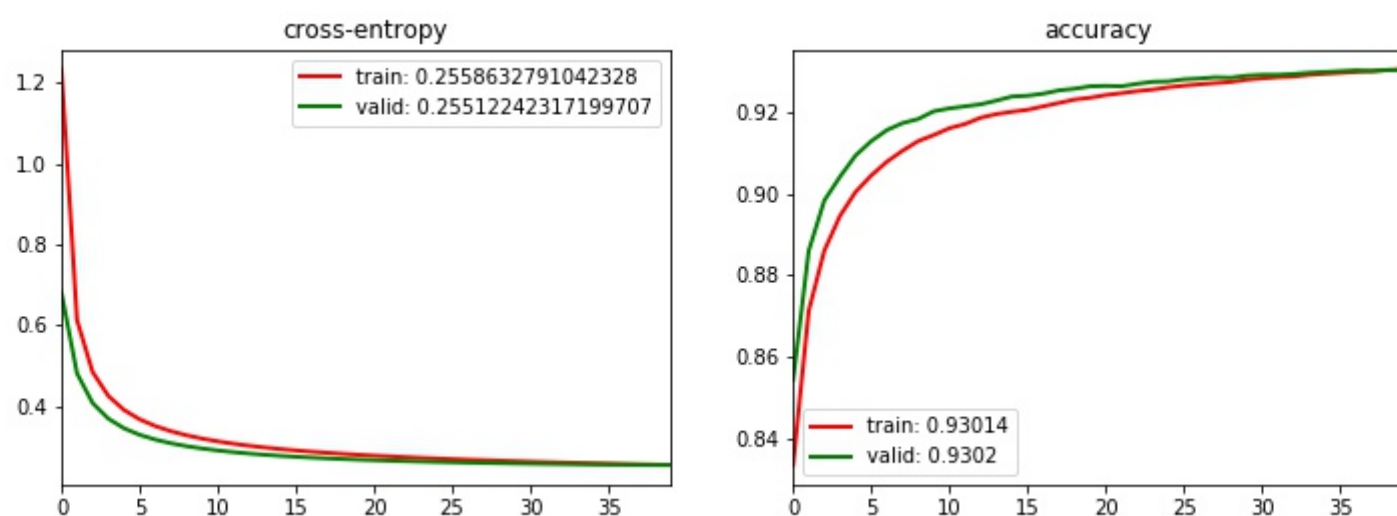
```
In [12]: 1 # Compute predictions
2 ##### YOUR CODE HERE ##### Logits for input_X, resulting shape should be [input_X.shape[0], 10]
3 logits = input_X @ W + b
4 # Logits = tf.add(tf.matmul(input_X, W), b)
5 ##### YOUR CODE HERE ##### apply tf.nn.softmax to Logits
6 probas = tf.nn.softmax(logits)
7 ##### YOUR CODE HERE ##### apply tf.argmax to find a class index with highest probability
8 classes = tf.argmax(probas, axis=1)
9
10 # Loss should be a scalar number: average loss over all the objects with tf.reduce_mean().
11 # Use tf.nn.softmax_cross_entropy_with_logits on top of one-hot encoded input_y and logits.
12 # It is identical to calculating cross-entropy on top of probas, but is more numerically friendly (read the docs).
13 ##### YOUR CODE HERE ##### cross-entropy loss
14 loss = tf.reduce_mean(tf.nn.softmax_cross_entropy_with_logits(labels=input_y, logits=logits))
15
16 # Use a default tf.train.AdamOptimizer to get an SGD step
17 ##### YOUR CODE HERE ##### optimizer step that minimizes the loss
18 step = tf.train.AdamOptimizer().minimize(loss)
```

WARNING:tensorflow:From <ipython-input-12-6e9b97a498be>:14: softmax_cross_entropy_with_logits (from tensorflow.python.ops.nn_ops) is deprecated and will be removed in a future version.
Instructions for updating:

Future major versions of TensorFlow will allow gradients to flow into the labels input on backprop by default.

See `tf.nn.softmax_cross_entropy_with_logits_v2`.

```
In [13]: 1 s.run(tf.global_variables_initializer())
2
3 BATCH_SIZE = 512
4 EPOCHS = 40
5
6 # for logging the progress right here in Jupyter (for those who don't have TensorBoard)
7 simpleTrainingCurves = matplotlib_utils.SimpleTrainingCurves("cross-entropy", "accuracy")
8
9 for epoch in range(EPOCHS): # we finish an epoch when we've looked at all training samples
10
11     batch_losses = []
12     for batch_start in range(0, X_train_flat.shape[0], BATCH_SIZE): # data is already shuffled
13         _, batch_loss = s.run([step, loss], {input_X: X_train_flat[batch_start:batch_start+BATCH_SIZE],
14                                             input_y: y_train_oh[batch_start:batch_start+BATCH_SIZE]})
15         # collect batch losses, this is almost free as we need a forward pass for backprop anyway
16         batch_losses.append(batch_loss)
17
18     train_loss = np.mean(batch_losses)
19     # print('train_loss', train_loss)
20     val_loss = s.run(loss, {input_X: X_val_flat, input_y: y_val_oh}) # this part is usually small
21     # print('val_loss', val_loss)
22     train_accuracy = accuracy_score(y_train, s.run(classes, {input_X: X_train_flat})) # this is slow and usually sk
23     valid_accuracy = accuracy_score(y_val, s.run(classes, {input_X: X_val_flat}))
24     simpleTrainingCurves.add(train_loss, val_loss, train_accuracy, valid_accuracy)
```



Submit a linear model

```
In [14]: 1 ## GRADED PART, DO NOT CHANGE!
2 # Testing shapes
3 grader.set_answer("9XaAS", grading_utils.get_tensors_shapes_string([W, b, input_X, input_y, logits, probas, classes])
4 # Validation Loss
5 grader.set_answer("vmogZ", s.run(loss, {input_X: X_val_flat, input_y: y_val_oh}))
6 # Validation accuracy
7 grader.set_answer("RMv95", accuracy_score(y_val, s.run(classes, {input_X: X_val_flat})))
```

```
In [15]: 1 # you can make submission with answers so far to check yourself at this stage
2 grader.submit(COURSERA_EMAIL, COURSERA_TOKEN)
```

Submitted to Coursera platform. See results on assignment page!

MLP with hidden layers

Previously we've coded a dense layer with matrix multiplication by hand. But this is not convenient, you have to create a lot of variables and your code becomes a mess. In TensorFlow there's an easier way to make a dense layer:

```
hidden1 = tf.layers.dense(inputs, 256, activation=tf.nn.sigmoid)
```

That will create all the necessary variables automatically. Here you can also choose an activation function (remember that we need it for a hidden layer!).

Now define the MLP with 2 hidden layers and restart training with the cell above.

You're aiming for ~0.97 validation accuracy here.

```
In [16]: 1 # write the code here to get a new `step` operation and then run the cell with training loop above.
2 # name your variables in the same way (e.g. logits, probas, classes, etc) for safety.
3 ### YOUR CODE HERE ###
4 hidden1 = tf.layers.dense(input_X, 256, activation=tf.nn.sigmoid)
5 hidden2 = tf.layers.dense(hidden1, 256, activation=tf.nn.sigmoid)
6 logits = tf.layers.dense(hidden2, 10)
7
8 probas = tf.nn.softmax(logits, 1)
9 classes = tf.argmax(probas, axis=1)
10 loss = tf.reduce_mean(tf.nn.softmax_cross_entropy_with_logits(labels=input_y, logits=logits))
11
12 step = tf.train.AdamOptimizer().minimize(loss)
```

WARNING:tensorflow:From <ipython-input-16-9f11ab2822ef>:4: dense (from tensorflow.python.layers.core) is deprecated and will be removed in a future version.

Instructions for updating:

Use keras.layers.dense instead.

WARNING:tensorflow:Entity <bound method Dense.call of <tensorflow.python.layers.core.Dense object at 0x1a773abe10>> could not be transformed and will be executed as-is. Please report this to the AutgoGraph team. When filing the bug, set the verbosity to 10 (on Linux, `export AUTOGRAPH_VERBOSITY=10`) and attach the full output. Cause: converting <bound method Dense.call of <tensorflow.python.layers.core.Dense object at 0x1a773abe10>>: AssertionError: Bad argument number for Name: 3, expecting 4

WARNING: Entity <bound method Dense.call of <tensorflow.python.layers.core.Dense object at 0x1a773abe10>> could not be transformed and will be executed as-is. Please report this to the AutgoGraph team. When filing the bug, set the verbosity to 10 (on Linux, `export AUTOGRAPH_VERBOSITY=10`) and attach the full output. Cause: converting <bound method Dense.call of <tensorflow.python.layers.core.Dense object at 0x1a773abe10>>: AssertionError: Bad argument number for Name: 3, expecting 4

WARNING:tensorflow:Entity <bound method Dense.call of <tensorflow.python.layers.core.Dense object at 0x1a773aba90>> could not be transformed and will be executed as-is. Please report this to the AutgoGraph team. When filing the bug, set the verbosity to 10 (on Linux, `export AUTOGRAPH_VERBOSITY=10`) and attach the full output. Cause: converting <bound method Dense.call of <tensorflow.python.layers.core.Dense object at 0x1a773aba90>>: AssertionError: Bad argument number for Name: 3, expecting 4

WARNING: Entity <bound method Dense.call of <tensorflow.python.layers.core.Dense object at 0x1a773aba90>> could not be transformed and will be executed as-is. Please report this to the AutgoGraph team. When filing the bug, set the verbosity to 10 (on Linux, `export AUTOGRAPH_VERBOSITY=10`) and attach the full output. Cause: converting <bound method Dense.call of <tensorflow.python.layers.core.Dense object at 0x1a773aba90>>: AssertionError: Bad argument number for Name: 3, expecting 4


WARNING:tensorflow:Entity <bound method Dense.call of <tensorflow.python.layers.core.Dense object at 0x1a77273410>> could not be transformed and will be executed as-is. Please report this to the AutgoGraph team. When filing the bug, set the verbosity to 10 (on Linux, `export AUTOGRAPH_VERBOSITY=10`) and attach the full output. Cause: converting <bound method Dense.call of <tensorflow.python.layers.core.Dense object at 0x1a77273410>>: AssertionError: Bad argument number for Name: 3, expecting 4

WARNING: Entity <bound method Dense.call of <tensorflow.python.layers.core.Dense object at 0x1a77273410>> could not be transformed and will be executed as-is. Please report this to the AutgoGraph team. When filing the bug, set the verbosity to 10 (on Linux, `export AUTOGRAPH_VERBOSITY=10`) and attach the full output. Cause: converting <bound method Dense.call of <tensorflow.python.layers.core.Dense object at 0x1a77273410>>: AssertionError: Bad argument number for Name: 3, expecting 4

```

In [17]: 1 s.run(tf.global_variables_initializer())
2
3 BATCH_SIZE = 512
4 EPOCHS = 40
5
6 # for logging the progress right here in Jupyter (for those who don't have TensorBoard)
7 simpleTrainingCurves = matplotlib_utils.SimpleTrainingCurves("cross-entropy", "accuracy")
8
9 for epoch in range(EPOCHS): # we finish an epoch when we've looked at all training samples
10
11     batch_losses = []
12     for batch_start in range(0, X_train_flat.shape[0], BATCH_SIZE): # data is already shuffled
13         _, batch_loss = s.run([step, loss], {input_X: X_train_flat[batch_start:batch_start+BATCH_SIZE],
14                                             input_y: y_train_oh[batch_start:batch_start+BATCH_SIZE]})
15         # collect batch losses, this is almost free as we need a forward pass for backprop anyway
16         batch_losses.append(batch_loss)
17
18     train_loss = np.mean(batch_losses)
19     # print('train_loss', train_loss)
20     val_loss = s.run(loss, {input_X: X_val_flat, input_y: y_val_oh}) # this part is usually small
21     # print('val_loss', val_loss)
22     train_accuracy = accuracy_score(y_train, s.run(classes, {input_X: X_train_flat})) # this is slow and usually sk
23     valid_accuracy = accuracy_score(y_val, s.run(classes, {input_X: X_val_flat}))
24     simpleTrainingCurves.add(train_loss, val_loss, train_accuracy, valid_accuracy)

```



Submit the MLP with 2 hidden layers

Run these cells after training the MLP with 2 hidden layers

```

In [18]: 1 ## GRADED PART, DO NOT CHANGE!
2 # Validation loss for MLP
3 grader.set_answer("i8bgs", s.run(loss, {input_X: X_val_flat, input_y: y_val_oh}))
4 # Validation accuracy for MLP
5 grader.set_answer("rE763", accuracy_score(y_val, s.run(classes, {input_X: X_val_flat})))

```

```

In [19]: 1 # you can make submission with answers so far to check yourself at this stage
2 grader.submit(COURSERA_EMAIL, COURSERA_TOKEN)

```

Submitted to Coursera platform. See results on assignment page!

```

In [ ]: 1

```

