


In [1]:

```
1 # set tf 1.x for colab
2 %tensorflow_version 1.x
```

UsageError: Line magic function `%tensorflow\_version` not found.

MNIST digits classification with TensorFlow



In [2]:

```
1 import numpy as np
2 from sklearn.metrics import accuracy_score
3 from matplotlib import pyplot as plt
4 %matplotlib inline
5 import tensorflow as tf
6 print("We're using TF", tf.__version__)
7
8 import sys
9 sys.path.append("../..")
10 import grading
11
12 import matplotlib_utils
13 from importlib import reload
14 reload(matplotlib_utils)
15
16 import grading_utils
17 reload(grading_utils)
18
19 import keras_utils
20 from keras_utils import reset_tf_session
```

/Users/Victor/anaconda3/envs/tfspark/lib/python3.7/site-packages/tensorflow/python/framework/dtypes.py:516: FutureWarning: Passing (type, 1) or 'ltype' as a synonym of type is deprecated; in a future version of numpy, it will be understood as (type, (1,)) / '(1,)type'.
 \_np\_qint8 = np.dtype [("qint8", np.int8, 1)]
/Users/Victor/anaconda3/envs/tfspark/lib/python3.7/site-packages/tensorflow/python/framework/dtypes.py:517: FutureWarning: Passing (type, 1) or 'ltype' as a synonym of type is deprecated; in a future version of numpy, it will be understood as (type, (1,)) / '(1,)type'.
 \_np\_quint8 = np.dtype [("quint8", np.uint8, 1)]
/Users/Victor/anaconda3/envs/tfspark/lib/python3.7/site-packages/tensorflow/python/framework/dtypes.py:518: FutureWarning: Passing (type, 1) or 'ltype' as a synonym of type is deprecated; in a future version of numpy, it will be understood as (type, (1,)) / '(1,)type'.
 \_np\_qint16 = np.dtype [("qint16", np.int16, 1)]
/Users/Victor/anaconda3/envs/tfspark/lib/python3.7/site-packages/tensorflow/python/framework/dtypes.py:519: FutureWarning: Passing (type, 1) or 'ltype' as a synonym of type is deprecated; in a future version of numpy, it will be understood as (type, (1,)) / '(1,)type'.
 \_np\_quint16 = np.dtype [("quint16", np.uint16, 1)]
/Users/Victor/anaconda3/envs/tfspark/lib/python3.7/site-packages/tensorflow/python/framework/dtypes.py:520: FutureWarning: Passing (type, 1) or 'ltype' as a synonym of type is deprecated; in a future version of numpy, it will be understood as (type, (1,)) / '(1,)type'.
 \_np\_qint32 = np.dtype [("qint32", np.int32, 1)]
/Users/Victor/anaconda3/envs/tfspark/lib/python3.7/site-packages/tensorflow/python/framework/dtypes.py:525: FutureWarning: Passing (type, 1) or 'ltype' as a synonym of type is deprecated; in a future version of numpy, it will be understood as (type, (1,)) / '(1,)type'.
 np\_resource = np.dtype [("resource", np.ubyte, 1)]

We're using TF 1.14.0

/Users/Victor/anaconda3/envs/tfspark/lib/python3.7/site-packages/tensorboard/compat/tensorflow\_stub/dtypes.py:541: FutureWarning: Passing (type, 1) or 'ltype' as a synonym of type is deprecated; in a future version of numpy, it will be understood as (type, (1,)) / '(1,)type'.
 \_np\_qint8 = np.dtype [("qint8", np.int8, 1)]
/Users/Victor/anaconda3/envs/tfspark/lib/python3.7/site-packages/tensorboard/compat/tensorflow\_stub/dtypes.py:542: FutureWarning: Passing (type, 1) or 'ltype' as a synonym of type is deprecated; in a future version of numpy, it will be understood as (type, (1,)) / '(1,)type'.
 \_np\_quint8 = np.dtype [("quint8", np.uint8, 1)]
/Users/Victor/anaconda3/envs/tfspark/lib/python3.7/site-packages/tensorboard/compat/tensorflow\_stub/dtypes.py:543: FutureWarning: Passing (type, 1) or 'ltype' as a synonym of type is deprecated; in a future version of numpy, it will be understood as (type, (1,)) / '(1,)type'.
 \_np\_qint16 = np.dtype [("qint16", np.int16, 1)]
/Users/Victor/anaconda3/envs/tfspark/lib/python3.7/site-packages/tensorboard/compat/tensorflow\_stub/dtypes.py:544: FutureWarning: Passing (type, 1) or 'ltype' as a synonym of type is deprecated; in a future version of numpy, it will be understood as (type, (1,)) / '(1,)type'.
 \_np\_quint16 = np.dtype [("quint16", np.uint16, 1)]
/Users/Victor/anaconda3/envs/tfspark/lib/python3.7/site-packages/tensorboard/compat/tensorflow\_stub/dtypes.py:545: FutureWarning: Passing (type, 1) or 'ltype' as a synonym of type is deprecated; in a future version of numpy, it will be understood as (type, (1,)) / '(1,)type'.
 \_np\_qint32 = np.dtype [("qint32", np.int32, 1)]
/Users/Victor/anaconda3/envs/tfspark/lib/python3.7/site-packages/tensorboard/compat/tensorflow\_stub/dtypes.py:550: FutureWarning: Passing (type, 1) or 'ltype' as a synonym of type is deprecated; in a future version of numpy, it will be understood as (type, (1,)) / '(1,)type'.
 np\_resource = np.dtype [("resource", np.ubyte, 1)]
Using TensorFlow backend.

Fill in your Coursera token and email

To successfully submit your answers to our grader, please fill in your Coursera submission token and email

```
In [3]: 1 grader = grading.Grader(assignment_key="XtD7ho3TEeiHQBLejjYAA",
2         all_parts=["9XaAS", "vmogZ", "RMv95", "i8bgs", "rE763"])
```

```
In [4]: 1 # token expires every 30 min
2 COURSERA_TOKEN = "xsKaGouklo4Fsk8K"
3 COURSERA_EMAIL = "lxwvictor@gmail.com"
```

## Look at the data


In this task we have 50000 28x28 images of digits from 0 to 9. We will train a classifier on this data.

```
In [5]: 1 import preprocessed_mnist
2 X_train, y_train, X_val, y_val, X_test, y_test = preprocessed_mnist.load_dataset()
```


```
In [6]: 1 # X contains rgb values divided by 255
2 print("X_train [shape %s] sample patch:\n" % (str(X_train.shape)), X_train[1, 15:20, 5:10])
3 print("A closeup of a sample patch:")
4 plt.imshow(X_train[1, 15:20, 5:10], cmap="Greys")
5 plt.show()
6 print("And the whole sample:")
7 plt.imshow(X_train[1], cmap="Greys")
8 plt.show()
9 print("y_train [shape %s] 10 samples:\n" % (str(y_train.shape)), y_train[:10])
```

X\_train [shape (50000, 28, 28)] sample patch:  
[[0. 0.29803922 0.96470588 0.98823529 0.43921569]  
[0. 0.33333333 0.98823529 0.90196078 0.09803922]  
[0. 0.33333333 0.98823529 0.8745098 0. ]  
[0. 0.33333333 0.98823529 0.56862745 0. ]  
[0. 0.3372549 0.99215686 0.88235294 0. ]]

A closeup of a sample patch:



And the whole sample:



y\_train [shape (50000,)] 10 samples:  
[5 0 4 1 9 2 1 3 1 4]

## Linear model

Your task is to train a linear classifier  $\vec{x} \rightarrow y$  with SGD using TensorFlow.

You will need to calculate a logit (a linear transformation)  $z_k$  for each class:

$$z_k = \vec{x} \cdot \vec{w}_k + b_k \quad k = 0..9$$

And transform logits  $z_k$  to valid probabilities  $p_k$  with softmax:

$$p_k = \frac{e^{z_k}}{\sum_{i=0}^9 e^{z_i}} \quad k = 0..9$$

We will use a cross-entropy loss to train our multi-class classifier:

$$\text{cross-entropy}(y, p) = - \sum_{k=0}^9 \log(p_k)[y = k]$$

where

$$[x] = \begin{cases} 1, & \text{if } x \text{ is true} \\ 0, & \text{otherwise} \end{cases}$$

Cross-entropy minimization pushes  $p_k$  close to 1 when  $y = k$ , which is what we want.

Here's the plan:

- Flatten the images (28x28 -> 784) with `X_train.reshape((X_train.shape[0], -1))` to simplify our linear model implementation
- Use a matrix placeholder for flattened `X_train`
- Convert `y_train` to one-hot encoded vectors that are needed for cross-entropy

- Use a shared variable  $w$  for all weights (a column  $\vec{w}_k$  per class) and  $b$  for all biases.
- Aim for ~0.93 validation accuracy

```
In [7]: 1 X_train_flat = X_train.reshape((X_train.shape[0], -1))
2 print(X_train_flat.shape)
3
4 X_val_flat = X_val.reshape((X_val.shape[0], -1))
5 print(X_val_flat.shape)
```

```
(50000, 784)
(10000, 784)
```

```
In [8]: 1 import keras
2
3 y_train_oh = keras.utils.to_categorical(y_train, 10)
4 y_val_oh = keras.utils.to_categorical(y_val, 10)
5
6 print(y_train_oh.shape)
7 print(y_train_oh[:3], y_train[:3])
```

```
(50000, 10)
[[0. 0. 0. 0. 0. 1. 0. 0. 0. 0.]
 [1. 0. 0. 0. 0. 0. 0. 0. 0. 0.]
 [0. 0. 0. 0. 1. 0. 0. 0. 0. 0.]] [5 0 4]
```

```
In [9]: 1 # run this again if you remake your graph
2 s = reset_tf_session()
```

WARNING:tensorflow:From ../../keras\_utils.py:68: The name tf.get\_default\_session is deprecated. Please use tf.compat.v1.get\_default\_session instead.

WARNING:tensorflow:From /Users/Victor/anaconda3/envs/tfspark/lib/python3.7/site-packages/keras/backend/tensorflow\_backend.py:95: The name tf.reset\_default\_graph is deprecated. Please use tf.compat.v1.reset\_default\_graph instead.

WARNING:tensorflow:From /Users/Victor/anaconda3/envs/tfspark/lib/python3.7/site-packages/keras/backend/tensorflow\_backend.py:98: The name tf.placeholder\_with\_default is deprecated. Please use tf.compat.v1.placeholder\_with\_default instead.

WARNING:tensorflow:From /Users/Victor/anaconda3/envs/tfspark/lib/python3.7/site-packages/keras/backend/tensorflow\_backend.py:102: The name tf.get\_default\_graph is deprecated. Please use tf.compat.v1.get\_default\_graph instead.

WARNING:tensorflow:From ../../keras\_utils.py:75: The name tf.ConfigProto is deprecated. Please use tf.compat.v1.ConfigProto instead.

```
In [10]: 1 # Model parameters: W and b
2 ### YOUR CODE HERE ### tf.get_variable(...) with shape[0] = 784
3 W = tf.get_variable('W', shape=(784, 10), dtype=tf.float32)
4 ### YOUR CODE HERE ### tf.get_variable(...)
5 b = tf.get_variable('b', shape=(10), dtype=tf.float32)
```

WARNING:tensorflow:From /Users/Victor/anaconda3/envs/tfspark/lib/python3.7/site-packages/tensorflow/python/ops/init\_ops.py:1251: calling VarianceScaling.\_\_init\_\_ (from tensorflow.python.ops.init\_ops) with dtype is deprecated and will be removed in a future version.

Instructions for updating:

Call initializer instance with the dtype argument instead of passing it to the constructor

```
In [11]: 1 # Placeholders for the input data
2 ### YOUR CODE HERE ### tf.placeholder(...) for flat X with shape[0] = None for any batch size
3 input_X = tf.placeholder(tf.float32, shape=(None, 784))
4 ### YOUR CODE HERE ### tf.placeholder(...) for one-hot encoded true labels
5 input_y = tf.placeholder(tf.int32, shape=(None, 10))
```

```
In [12]: 1 # Compute predictions
2 ### YOUR CODE HERE ### logits for input_X, resulting shape should be [input_X.shape[0], 10]
3 logits = input_X @ W + b
4 # logits = tf.add(tf.matmul(input_X, W), b)
5 ### YOUR CODE HERE ### apply tf.nn.softmax to logits
6 probas = tf.nn.softmax(logits)
7 ### YOUR CODE HERE ### apply tf.argmax to find a class index with highest probability
8 classes = tf.argmax(probas, axis=1)
9
10 # Loss should be a scalar number: average loss over all the objects with tf.reduce_mean().
11 # Use tf.nn.softmax_cross_entropy_with_logits on top of one-hot encoded input_y and logits.
12 # It is identical to calculating cross-entropy on top of probas, but is more numerically friendly (read the docs).
13 ### YOUR CODE HERE ### cross-entropy loss
14 loss = tf.reduce_mean(tf.nn.softmax_cross_entropy_with_logits(labels=input_y, logits=logits))
15
16 # Use a default tf.train.AdamOptimizer to get an SGD step
17 ### YOUR CODE HERE ### optimizer step that minimizes the loss
18 step = tf.train.AdamOptimizer().minimize(loss)
```

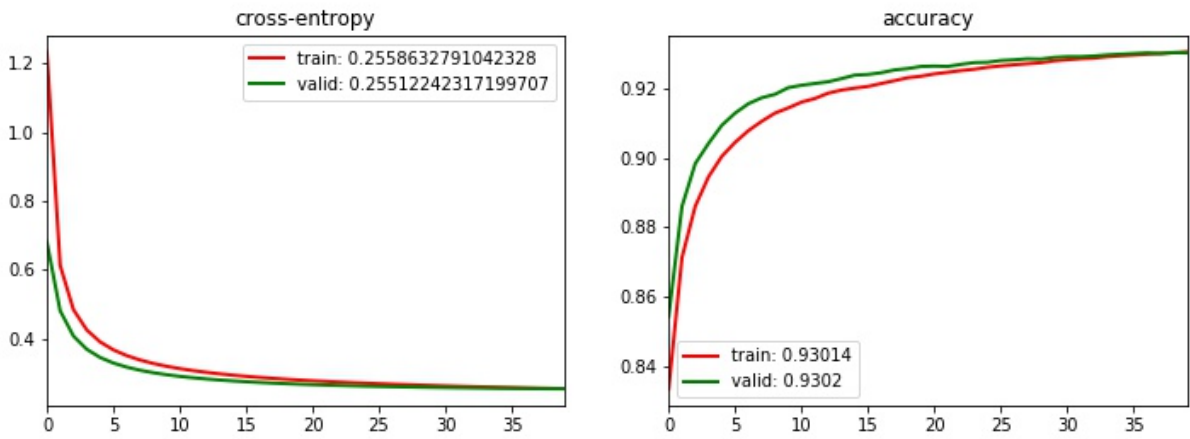
WARNING:tensorflow:From <ipython-input-12-6e9b97a498be>:14: softmax\_cross\_entropy\_with\_logits (from tensorflow.python.ops.nn\_ops) is deprecated and will be removed in a future version.

Instructions for updating:

Future major versions of TensorFlow will allow gradients to flow into the labels input on backprop by default.

See `tf.nn.softmax\_cross\_entropy\_with\_logits\_v2`.

```
In [13]: 1 s.run(tf.global_variables_initializer())
2
3 BATCH_SIZE = 512
4 EPOCHS = 40
5
6 # for logging the progress right here in Jupyter (for those who don't have TensorBoard)
7 simpleTrainingCurves = matplotlib_utils.SimpleTrainingCurves("cross-entropy", "accuracy")
8
9 for epoch in range(EPOCHS): # we finish an epoch when we've looked at all training samples
10
11     batch_losses = []
12     for batch_start in range(0, X_train_flat.shape[0], BATCH_SIZE): # data is already shuffled
13         _, batch_loss = s.run([step, loss], {input_X: X_train_flat[batch_start:batch_start+BATCH_SIZE],
14                                             input_y: y_train_oh[batch_start:batch_start+BATCH_SIZE]})
15         # collect batch losses, this is almost free as we need a forward pass for backprop anyway
16         batch_losses.append(batch_loss)
17
18     train_loss = np.mean(batch_losses)
19     # print('train_loss', train_loss)
20     val_loss = s.run(loss, {input_X: X_val_flat, input_y: y_val_oh}) # this part is usually small
21     # print('val_loss', val_loss)
22     train_accuracy = accuracy_score(y_train, s.run(classes, {input_X: X_train_flat})) # this is slow and usually skipped
23     valid_accuracy = accuracy_score(y_val, s.run(classes, {input_X: X_val_flat}))
24     simpleTrainingCurves.add(train_loss, val_loss, train_accuracy, valid_accuracy)
```



## Submit a linear model

```
In [14]: 1 ## GRADED PART, DO NOT CHANGE!
2 # Testing shapes
3 grader.set_answer("9XaAS", grading_utils.get_tensors_shapes_string([W, b, input_X, input_y, logits, probas, classes]))
4 # Validation loss
5 grader.set_answer("vmogZ", s.run(loss, {input_X: X_val_flat, input_y: y_val_oh}))
6 # Validation accuracy
7 grader.set_answer("RMv95", accuracy_score(y_val, s.run(classes, {input_X: X_val_flat})))
```

```
In [15]: 1 # you can make submission with answers so far to check yourself at this stage
2 grader.submit(COURSERA_EMAIL, COURSERA_TOKEN)
```

Submitted to Coursera platform. See results on assignment page!

## MLP with hidden layers

Previously we've coded a dense layer with matrix multiplication by hand. But this is not convenient, you have to create a lot of variables and your code becomes a mess. In TensorFlow there's an easier way to make a dense layer:

```
hidden1 = tf.layers.dense(inputs, 256, activation=tf.nn.sigmoid)
```

That will create all the necessary variables automatically. Here you can also choose an activation function (remember that we need it for a hidden layer!).

Now define the MLP with 2 hidden layers and restart training with the cell above.

You're aiming for ~0.97 validation accuracy here.



```
In [16]: 1 # write the code here to get a new `step` operation and then run the cell with training loop above.
2 # name your variables in the same way (e.g. logits, probas, classes, etc) for safety.
3 ### YOUR CODE HERE ###
4 hidden1 = tf.layers.dense(input_X, 256, activation=tf.nn.sigmoid)
5 hidden2 = tf.layers.dense(hidden1, 256, activation=tf.nn.sigmoid)
6 logits = tf.layers.dense(hidden2, 10)
7
8 probas = tf.nn.softmax(logits, 1)
9 classes = tf.argmax(probas, axis=1)
10 loss = tf.reduce_mean(tf.nn.softmax_cross_entropy_with_logits(labels=input_y, logits=logits))
11
12 step = tf.train.AdamOptimizer().minimize(loss)
```

WARNING:tensorflow:From <ipython-input-16-9f11ab2822ef>:4: dense (from tensorflow.python.layers.core) is deprecated and will be removed in a future version.

Instructions for updating:

Use keras.layers.dense instead.

WARNING:tensorflow:Entity <bound method Dense.call of <tensorflow.python.layers.core.Dense object at 0x1a773abel0>> could not be transformed and will be executed as-is. Please report this to the AutgoGraph team. When filing the bug, set the verbosity to 10 (on Linux, `export AUTOGRAPH\_VERBOSITY=10`) and attach the full output. Cause: converting <bound method Dense.call of <tensorflow.python.layers.core.Dense object at 0x1a773abel0>>: AssertionError: Bad argument number for Name: 3, expecting 4

WARNING: Entity <bound method Dense.call of <tensorflow.python.layers.core.Dense object at 0x1a773abel0>> could not be transformed and will be executed as-is. Please report this to the AutgoGraph team. When filing the bug, set the verbosity to 10 (on Linux, `export AUTOGRAPH\_VERBOSITY=10`) and attach the full output. Cause: converting <bound method Dense.call of <tensorflow.python.layers.core.Dense object at 0x1a773abel0>>: AssertionError: Bad argument number for Name: 3, expecting 4


WARNING:tensorflow:Entity <bound method Dense.call of <tensorflow.python.layers.core.Dense object at 0x1a773aba90>> could not be transformed and will be executed as-is. Please report this to the AutgoGraph team. When filing the bug, set the verbosity to 10 (on Linux, `export AUTOGRAPH\_VERBOSITY=10`) and attach the full output. Cause: converting <bound method Dense.call of <tensorflow.python.layers.core.Dense object at 0x1a773aba90>>: AssertionError: Bad argument number for Name: 3, expecting 4

WARNING: Entity <bound method Dense.call of <tensorflow.python.layers.core.Dense object at 0x1a773aba90>> could not be transformed and will be executed as-is. Please report this to the AutgoGraph team. When filing the bug, set the verbosity to 10 (on Linux, `export AUTOGRAPH\_VERBOSITY=10`) and attach the full output. Cause: converting <bound method Dense.call of <tensorflow.python.layers.core.Dense object at 0x1a773aba90>>: AssertionError: Bad argument number for Name: 3, expecting 4

WARNING:tensorflow:Entity <bound method Dense.call of <tensorflow.python.layers.core.Dense object at 0x1a77273410>> could not be transformed and will be executed as-is. Please report this to the AutgoGraph team. When filing the bug, set the verbosity to 10 (on Linux, `export AUTOGRAPH\_VERBOSITY=10`) and attach the full output. Cause: converting <bound method Dense.call of <tensorflow.python.layers.core.Dense object at 0x1a77273410>>: AssertionError: Bad argument number for Name: 3, expecting 4

WARNING: Entity <bound method Dense.call of <tensorflow.python.layers.core.Dense object at 0x1a77273410>> could not be transformed and will be executed as-is. Please report this to the AutgoGraph team. When filing the bug, set the verbosity to 10 (on Linux, `export AUTOGRAPH\_VERBOSITY=10`) and attach the full output. Cause: converting <bound method Dense.call of <tensorflow.python.layers.core.Dense object at 0x1a77273410>>: AssertionError: Bad argument number for Name: 3, expecting 4

```
In [17]: 1 s.run(tf.global_variables_initializer())
2
3 BATCH_SIZE = 512
4 EPOCHS = 40
5
6 # for logging the progress right here in Jupyter (for those who don't have TensorBoard)
7 simpleTrainingCurves = matplotlib_utils.SimpleTrainingCurves("cross-entropy", "accuracy")
8
9 for epoch in range(EPOCHS): # we finish an epoch when we've looked at all training samples
10
11     batch_losses = []
12     for batch_start in range(0, X_train_flat.shape[0], BATCH_SIZE): # data is already shuffled
13         _, batch_loss = s.run([step, loss], {input_X: X_train_flat[batch_start:batch_start+BATCH_SIZE],
14                                             input_y: y_train_oh[batch_start:batch_start+BATCH_SIZE]})
15         # collect batch losses, this is almost free as we need a forward pass for backprop anyway
16         batch_losses.append(batch_loss)
17
18     train_loss = np.mean(batch_losses)
19     # print('train_loss', train_loss)
20     val_loss = s.run(loss, {input_X: X_val_flat, input_y: y_val_oh}) # this part is usually small
21     # print('val_loss', val_loss)
22     train_accuracy = accuracy_score(y_train, s.run(classes, {input_X: X_train_flat})) # this is slow and usually skipped
23     valid_accuracy = accuracy_score(y_val, s.run(classes, {input_X: X_val_flat}))
24     simpleTrainingCurves.add(train_loss, val_loss, train_accuracy, valid_accuracy)
```



## Submit the MLP with 2 hidden layers

Run these cells after training the MLP with 2 hidden layers

```
In [18]: 1 ## GRADED PART, DO NOT CHANGE!
2 # Validation loss for MLP
3 grader.set_answer("i8bgs", s.run(loss, {input_X: X_val_flat, input_y: y_val_oh}))
4 # Validation accuracy for MLP
5 grader.set_answer("rE763", accuracy_score(y_val, s.run(classes, {input_X: X_val_flat})))
```

```
In [19]: 1 # you can make submission with answers so far to check yourself at this stage
2 grader.submit(COURSERA_EMAIL, COURSERA_TOKEN)
```

Submitted to Coursera platform. See results on assignment page!

In [ ]:

1