

```
In [1]: 1 # set tf 1.x for colab
2 %tensorflow_version 1.x
```

UsageError: Line magic function `%tensorflow_version` not found.

```
In [2]: 1 import warnings
2 warnings.filterwarnings('ignore', category=DeprecationWarning)
3 warnings.filterwarnings('ignore', category=FutureWarning)
```

Generating human faces with Adversarial Networks



© research.nvidia.com

This time we'll train a neural net to generate plausible human faces in all their subtlety: appearance, expression, accessories, etc. 'Cuz when us machines gonna take over Earth, there won't be any more faces left. We want to preserve this data for future iterations. Yikes...

Based on <https://github.com/Lasagne/Recipes/pull/94> (<https://github.com/Lasagne/Recipes/pull/94>).

```
In [3]: 1 import sys
2 sys.path.append("..")
3 import grading
4 import download_utils
5 import tqdm_utils
```

```
In [4]: 1 download_utils.link_week_4_resources()
```

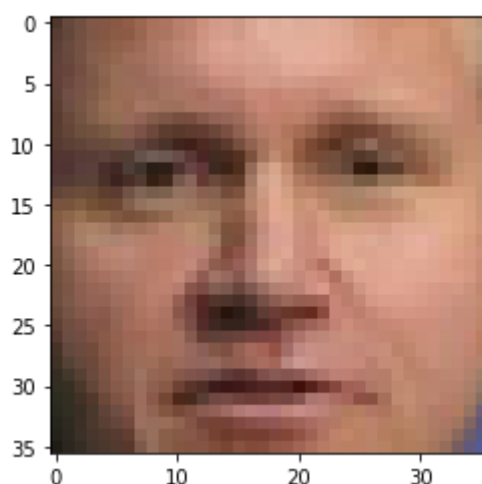
```
In [5]: 1 import matplotlib.pyplot as plt
2 %matplotlib inline
3 import numpy as np
4 plt.rcParams.update({'axes.titlesize': 'small'})
5
6 from sklearn.datasets import load_digits
7 #The following line fetches you two datasets: images, usable for autoencoder training and attributes.
8 #Those attributes will be required for the final part of the assignment (applying smiles), so please keep them in mi
9 from lfw_dataset import load_lfw_dataset
10 data, attrs = load_lfw_dataset(dimx=36, dimy=36)
11
12 #preprocess faces
13 data = np.float32(data)/255.
14
15 IMG_SHAPE = data.shape[1:]
```

100%


13233/13233 [00:30<00:00, 438.69it/s]

```
In [6]: 1 #print random image
2 plt.imshow(data[np.random.randint(data.shape[0])], cmap="gray", interpolation="none")
```

Out[6]: <matplotlib.image.AxesImage at 0x2164cdf8b48>



Generative adversarial nets 101



@ torch.github.io

Deep learning is simple, isn't it?

- build some network that generates the face (small image)
- make up a **measure of how good that face is**
- optimize with gradient descent :)

The only problem is: how can we engineers tell well-generated faces from bad? And i bet you we won't ask a designer for help.

If we can't tell good faces from bad, we delegate it to yet another neural network!

That makes the two of them:

- **Generator** - takes random noise for inspiration and tries to generate a face sample.
 - Let's call him **G(z)**, where z is a gaussian noise.
- **Discriminator** - takes a face sample and tries to tell if it's great or fake.
 - Predicts the probability of input image being a **real face**
 - Let's call him **D(x)**, x being an image.
 - **D(x)** is a prediction for real image and **D(G(z))** is prediction for the face made by generator.

Before we dive into training them, let's construct the two networks.

```
In [7]: 1 import tensorflow as tf
2 from keras_utils import reset_tf_session
3 s = reset_tf_session()
4
5 import keras
6 from keras.models import Sequential
7 from keras import layers as L
```

WARNING:tensorflow:From ..\keras_utils.py:68: The name tf.get_default_session is deprecated. Please use tf.compat.v1.get_default_session instead.

WARNING:tensorflow:From ..\keras_utils.py:75: The name tf.ConfigProto is deprecated. Please use tf.compat.v1.ConfigProto instead.

WARNING:tensorflow:From ..\keras_utils.py:77: The name tf.InteractiveSession is deprecated. Please use tf.compat.v1.InteractiveSession instead.

Using TensorFlow backend.

```
In [8]: 1 CODE_SIZE = 256
2
3 generator = Sequential()
4 generator.add(L.InputLayer([CODE_SIZE],name='noise'))
5 generator.add(L.Dense(10*8*8, activation='elu'))
6
7 generator.add(L.Reshape((8,8,10)))
8 generator.add(L.Deconvolution2D(64,kernel_size=(5,5),activation='elu'))
9 generator.add(L.Deconvolution2D(64,kernel_size=(5,5),activation='elu'))
10 generator.add(L.UpSampling2D(size=(2,2)))
11 generator.add(L.Deconvolution2D(32,kernel_size=3,activation='elu'))
12 generator.add(L.Deconvolution2D(32,kernel_size=3,activation='elu'))
13 generator.add(L.Deconvolution2D(32,kernel_size=3,activation='elu'))
14
15 generator.add(L.Conv2D(3,kernel_size=3,activation=None))
16
```

```
In [9]: 1 assert generator.output_shape[1:] == IMG_SHAPE, "generator must output an image of shape %s, but instead it produces %s"
```

Discriminator

- Discriminator is your usual convolutional network with interlooping convolution and pooling layers
- The network does not include dropout/batchnorm to avoid learning complications.
- We also regularize the pre-output layer to prevent discriminator from being too certain.

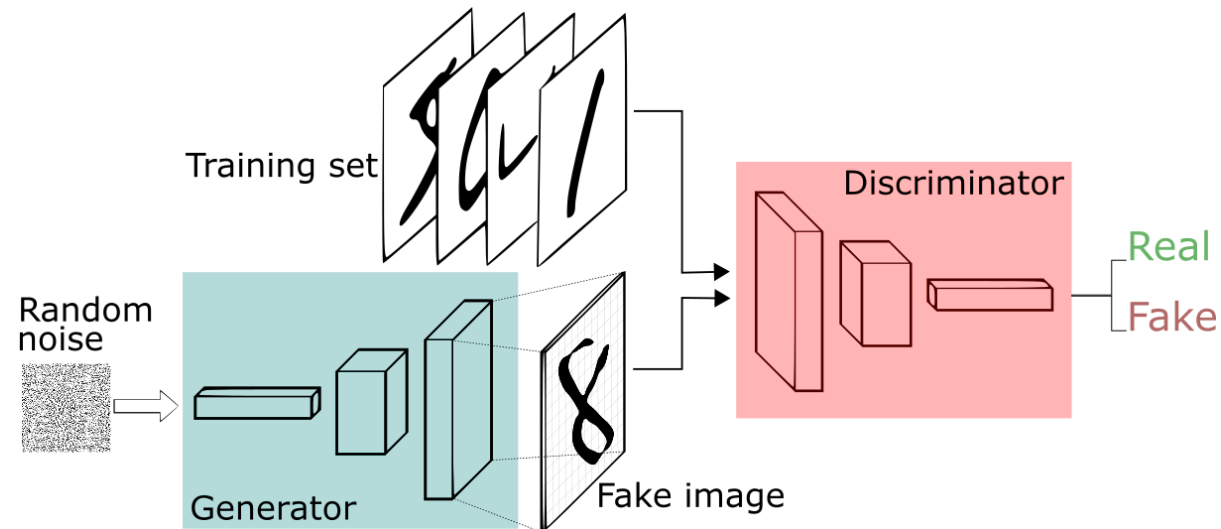
```
In [10]: 1 discriminator = Sequential()
2
3 discriminator.add(L.InputLayer(IMG_SHAPE))
4
5 # <build discriminator body>
6 discriminator.add(L.Conv2D(16,[2,2],padding='same',activation='relu'))
7 discriminator.add(L.Conv2D(32,[2,2],padding='same',activation='relu'))
8 discriminator.add(L.MaxPool2D())
9
10 discriminator.add(L.Conv2D(64,[2,2],padding='same',activation='relu'))
11 discriminator.add(L.Conv2D(128,[2,2],padding='same',activation='relu'))
12 discriminator.add(L.MaxPool2D())
13
14
15 discriminator.add(L.Flatten())
16 discriminator.add(L.Dense(256,activation='tanh'))
17 discriminator.add(L.Dense(2,activation=tf.nn.log_softmax))
```

WARNING:tensorflow:From C:\Users\Xiaowei\Anaconda3\envs\tfspark\lib\site-packages\keras\backend\tensorflow_backend.py:4070: The name tf.nn.max_pool is deprecated. Please use tf.nn.max_pool2d instead.

Training

We train the two networks concurrently:

- Train **discriminator** to better distinguish real data from **current** generator
- Train **generator** to make discriminator think generator is real
- Since discriminator is a differentiable neural network, we train both with gradient descent.



© deeplearning4j.org

Training is done iteratively until discriminator is no longer able to find the difference (or until you run out of patience).

Tricks:

- Regularize discriminator output weights to prevent explosion
- Train generator with **adam** to speed up training. Discriminator trains with SGD to avoid problems with momentum.
- More: <https://github.com/soumith/ganhacks> (<https://github.com/soumith/ganhacks>)

```
In [11]: 1 noise = tf.placeholder('float32',[None,CODE_SIZE])
2 real_data = tf.placeholder('float32',[None,]+list(IMG_SHAPE))
3
4 logp_real = discriminator(real_data)
5
6 generated_data = generator(noise) #<gen(noise)>
7
8 logp_gen = discriminator(generated_data) #<log P(real | gen(noise))>
```

```
In [12]: 1 #####
2 #discriminator training#
3 #####
4
5 d_loss = -tf.reduce_mean(logp_real[:,1] + logp_gen[:,0])
6
7 #regularize
8 d_loss += tf.reduce_mean(discriminator.layers[-1].kernel**2)
9
10 #optimize
11 disc_optimizer = tf.train.GradientDescentOptimizer(1e-3).minimize(d_loss,var_list=discriminator.trainable_weights)
```

WARNING:tensorflow:From C:\Users\Xiaowei\Anaconda3\envs\tfspark\lib\site-packages\tensorflow\python\ops\math_grad.py:1205: add_dispatch_support.<locals>.wrapper (from tensorflow.python.ops.array_ops) is deprecated and will be removed in a future version.

Instructions for updating:

Use tf.where in 2.0, which has the same broadcast rule as np.where

```
In [13]: 1 #####
2 ###generator training###
3 #####
4
5 g_loss = -tf.reduce_mean(logp_gen[:,1]) #<generator loss>
6
7 gen_optimizer = tf.train.AdamOptimizer(1e-4).minimize(g_loss,var_list=generator.trainable_weights)
```

```
In [14]: 1 s.run(tf.global_variables_initializer())
```

Auxiliary functions

Here we define a few helper functions that draw current data distributions and sample training batches.

```
In [15]: 1 def sample_noise_batch(bsize):
2     return np.random.normal(size=(bsize, CODE_SIZE)).astype('float32')
3
4 def sample_data_batch(bsize):
5     idxs = np.random.choice(np.arange(data.shape[0]), size=bsize)
6     return data[idxs]
7
8 def sample_images(nrow,ncol, sharp=False):
9     images = generator.predict(sample_noise_batch(bsize=nrow*ncol))
10    if np.var(images)!=0:
11        images = images.clip(np.min(data),np.max(data))
12    for i in range(nrow*ncol):
13        plt.subplot(nrow,ncol,i+1)
14        if sharp:
15            plt.imshow(images[i].reshape(IMG_SHAPE),cmap="gray", interpolation="none")
16        else:
17            plt.imshow(images[i].reshape(IMG_SHAPE),cmap="gray")
18    plt.show()
19
20 def sample_probas(bsize):
21    plt.title('Generated vs real data')
22    plt.hist(np.exp(discriminator.predict(sample_data_batch(bsize)))[:,1],
23            label='D(x)', alpha=0.5,range=[0,1])
24    plt.hist(np.exp(discriminator.predict(generator.predict(sample_noise_batch(bsize))))[:,1],
25            label='D(G(z))',alpha=0.5,range=[0,1])
26    plt.legend(loc='best')
27    plt.show()
```

Training

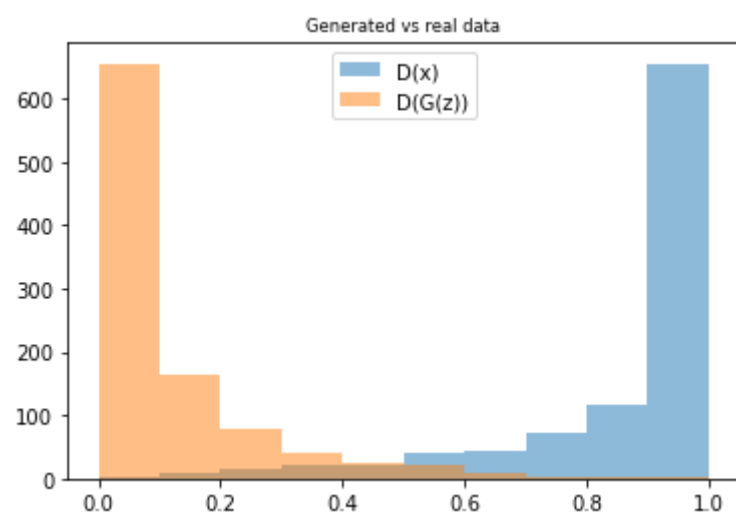

Main loop. We just train generator and discriminator in a loop and plot results once every N iterations.

```

In [16]: 1 %%time
2 from IPython import display
3
4 for epoch in tqdm_utils.tqdm_notebook_failsafe(range(15000)):
5
6     feed_dict = {
7         real_data:sample_data_batch(100),
8         noise:sample_noise_batch(100)
9     }
10
11     for i in range(5):
12         s.run(disc_optimizer,feed_dict)
13
14     s.run(gen_optimizer,feed_dict)
15
16     if epoch %100==0:
17         display.clear_output(wait=True)
18         print('Epoch status:', epoch)
19         sample_images(2,3,True)
20         sample_probas(1000)
21

```

Epoch status: 14900



Wall time: 51min 24s


```

In [ ]: 1 from submit_honor import submit_honor
2 submit_honor((generator, discriminator), <YOUR_EMAIL>, <YOUR_TOKEN>)

```



```
In [18]: 1 #The network was trained for about 15k iterations.
2 #Training for Longer yields MUCH better results
3 plt.figure(figsize=[16,24])
4 sample_images(16,8)
```



In []: 1