

```
In [1]: 1 # set tf 1.x for colab
        2 %tensorflow_version 1.x
```

UsageError: Line magic function `%tensorflow_version` not found.

```
In [2]: 1 import warnings
        2 warnings.filterwarnings('ignore', category=DeprecationWarning)
        3 warnings.filterwarnings('ignore', category=FutureWarning)
```

Generating names with recurrent neural networks

This time you'll find yourself delving into the heart (and other intestines) of recurrent neural networks on a class of toy problems.

Struggle to find a name for the variable? Let's see how you'll come up with a name for your son/daughter. Surely no human has expertise over what is a good child name, so let us train RNN instead;

It's dangerous to go alone, take these:

```
In [3]: 1 import tensorflow as tf
        2 print(tf.__version__)
        3 import numpy as np
        4 import matplotlib.pyplot as plt
        5 %matplotlib inline
        6 import os
        7 import sys
        8 sys.path.append("..")
        9 import keras_utils
       10 import tqdm_utils
```

1.14.0

Using TensorFlow backend.

Load data

The dataset contains ~8k earthling names from different cultures, all in latin transcript.

This notebook has been designed so as to allow you to quickly swap names for something similar: deep learning article titles, IKEA furniture, pokemon names, etc.

```
In [4]: 1 start_token = " " # so that the network knows that we're generating a first token
        2
        3 # this is the token for padding,
        4 # we will add fake pad token at the end of names
        5 # to make them of equal size for further batching
        6 pad_token = "#"
        7
        8 with open("names") as f:
        9     names = f.read()[:-1].split('\n')
       10     names = [start_token + name for name in names]
```


```
In [5]: 1 print('number of samples:', len(names))
        2 for x in names[:1000]:
        3     print(x)
```

number of samples: 7944

Abagael
Claresta
Glory
Liliane
Prissie
Geeta
Giovanne
Piggy

```
In [6]: 1 MAX_LENGTH = max(map(len, names))
2 print("max length:", MAX_LENGTH)
3
4 plt.title('Sequence length distribution')
5 plt.hist(list(map(len, names)), bins=25);
```

max length: 16



Text processing

First we need to collect a "vocabulary" of all unique tokens i.e. unique characters. We can then encode inputs as a sequence of character ids.

```
In [7]: 1 ### YOUR CODE HERE: all unique characters go here, padding included!
2 tokens = sorted(set(''.join(names + [start_token, pad_token])))
3
4 tokens = list(tokens)
5 n_tokens = len(tokens)
6 print ('n_tokens:', n_tokens)
7
8 assert 50 < n_tokens < 60
```

n_tokens: 56

Cast everything from symbols into identifiers

Tensorflow string manipulation is a bit tricky, so we'll work around it. We'll feed our recurrent neural network with ids of characters from our dictionary.

To create such dictionary, let's assign token_to_id

```
In [8]: 1 ### YOUR CODE HERE: create a dictionary of {symbol -> its index in tokens}
2 token_to_id = {v:k for k, v in enumerate(tokens)}
3
4 assert len(tokens) == len(token_to_id), "dictionaries must have same size"
```

1	token_to_id
---	-------------

```
Out[9]: {' ': 0,
         '#': 1,
         '"': 2,
         '-': 3,
         'A': 4,
         'B': 5,
         'C': 6,
         'D': 7,
         'E': 8,
         'F': 9,
         'G': 10,
         'H': 11,
         'I': 12,
         'J': 13,
         'K': 14,
         'L': 15,
         'M': 16,
         'N': 17,
         'O': 18,
         'P': 19,
```

```

1 def to_matrix(names, max_len=None, pad=token_to_id[pad_token], dtype=np.int32):
2     """Casts a list of names into rnn-digestable padded matrix"""
3
4     max_len = max_len or max(map(len, names))
5     names_ix = np.zeros([len(names), max_len], dtype) + pad
6
7     for i in range(len(names)):
8         name_ix = list(map(token_to_id.get, names[i]))
9         names_ix[i, :len(name_ix)] = name_ix
10
11     return names_ix

```

```
1 # Example: cast 4 random names to padded matrices (so that we can easily batch them)
2 print('\n'.join(names[:2000]))
3 print(to_matrix(names[:2000]))
```

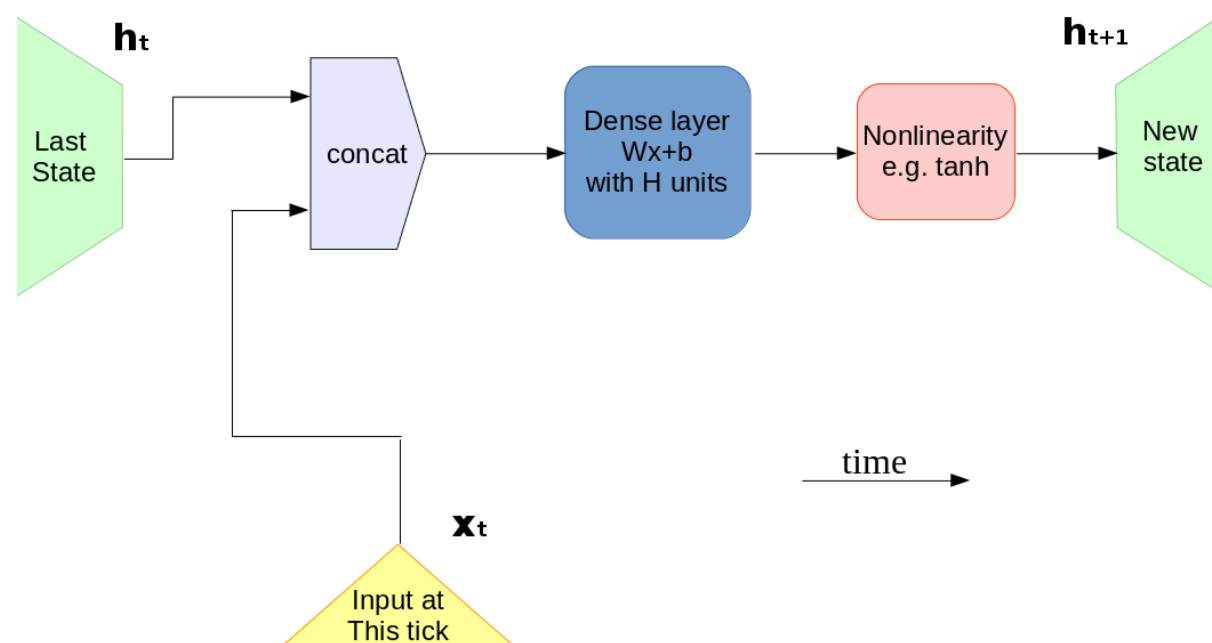
```

Abagael
Glory
Prissie
Giovanne
[[ 0  4 31 ...  1  1  1]
 [ 0  4 31 ...  1  1  1]
 [ 0  4 31 ...  1  1  1]
 ...
 [ 0 10 41 ...  1  1  1]
 [ 0 10 41 ...  1  1  1]
 [ 0 10 41 ...  1  1  1]]

```

Defining a recurrent neural network

We can rewrite recurrent neural network as a consecutive application of dense layer to input x_t and previous rnn state h_t . This is exactly what we're gonna do now.



Since we're training a language model, there should also be:

- An embedding layer that converts character id x_t to a vector.
- An output layer that predicts probabilities of next phoneme based on h_{t+1}

```
In [12]: 1 # remember to reset your session if you change your graph!
2 s = keras_utils.reset_tf_session()
```


WARNING:tensorflow:From ..\keras_utils.py:68: The name tf.get_default_session is deprecated. Please use tf.compat.v1.get_default_session instead.

WARNING:tensorflow:From ..\keras_utils.py:75: The name tf.ConfigProto is deprecated. Please use tf.compat.v1.ConfigProto instead.

WARNING:tensorflow:From ..\keras_utils.py:77: The name tf.InteractiveSession is deprecated. Please use tf.compat.v1.InteractiveSession instead.

```
In [13]: 1 import keras
2 from keras.layers import concatenate, Dense, Embedding
3
4 rnn_num_units = 64 # size of hidden state
5 embedding_size = 16 # for characters
6
7 # Let's create layers for our recurrent network
8 # Note: we create layers but we don't "apply" them yet (this is a "functional API" of Keras)
9 # Note: set the correct activation (from keras.activations) to Dense layers!
10
11 # an embedding layer that converts character ids into embeddings
12 embed_x = Embedding(n_tokens, embedding_size)
13
14 # a dense layer that maps input and previous state to new hidden state, [x_t, h_t] -> h_{t+1}
15 ### YOUR CODE HERE
16 get_h_next = Dense(rnn_num_units, activation='tanh')
17
18 # a dense layer that maps current hidden state to probabilities of characters [h_{t+1}] -> P(x_{t+1}|h_{t+1})
19 ### YOUR CODE HERE
20 get_probas = Dense(n_tokens, activation='softmax')
```

We will generate names character by character starting with `start_token` :



```
In [14]: 1 def rnn_one_step(x_t, h_t):
2     """
3     Recurrent neural network step that produces
4     probabilities for next token x_{t+1} and next state h_{t+1}
5     given current input x_t and previous state h_t.
6     We'll call this method repeatedly to produce the whole sequence.
7
8     You're supposed to "apply" above layers to produce new tensors.
9     Follow inline instructions to complete the function.
10    """
11    # convert character id into embedding
12    x_t_emb = embed_x(tf.reshape(x_t, [-1, 1]))[:, 0]
13
14    # concatenate x_t embedding and previous h_t state
15    ### YOUR CODE HERE
16    x_and_h = concatenate([x_t_emb, h_t], axis=-1)
17
18    # compute next state given x_and_h
19    ### YOUR CODE HERE
20    h_next = get_h_next(x_and_h)
21
22    # get probabilities for language model P(x_{next}|h_{next})
23    ### YOUR CODE HERE
24    output_probas = get_probas(h_next)
25
26    return output_probas, h_next
```

RNN: loop

Once `rnn_one_step` is ready, let's apply it in a loop over name characters to get predictions.

Let's assume that all names are at most length-16 for now, so we can simply iterate over them in a for loop.

```
In [15]: 1 input_sequence = tf.placeholder(tf.int32, (None, MAX_LENGTH)) # batch of token ids
2 batch_size = tf.shape(input_sequence)[0]
3
4 predicted_probas = []
5 h_prev = tf.zeros([batch_size, rnn_num_units]) # initial hidden state
6
7 for t in range(MAX_LENGTH):
8     x_t = input_sequence[:, t] # column t
9     probas_next, h_next = rnn_one_step(x_t, h_prev)
10
11     h_prev = h_next
12     predicted_probas.append(probas_next)
13
14 # combine predicted_probas into [batch, time, n_tokens] tensor
15 predicted_probas = tf.transpose(tf.stack(predicted_probas), [1, 0, 2])
16
17 # next to last token prediction is not needed
18 predicted_probas = predicted_probas[:, :-1, :]
```

RNN: loss and gradients

Let's gather a matrix of predictions for $P(x_{next}|h)$ and the corresponding correct answers.

We will flatten our matrices to shape [None, n_tokens] to make it easier.

Our network can then be trained by minimizing crossentropy between predicted probabilities and those answers.

```
In [16]: 1 # flatten predictions to [batch*time, n_tokens]
2 predictions_matrix = tf.reshape(predicted_probas, [-1, n_tokens])
3
4 # flatten answers (next tokens) and one-hot encode them
5 answers_matrix = tf.one_hot(tf.reshape(input_sequence[:, 1:], [-1]), n_tokens)
```

Usually it's a good idea to ignore gradients of loss for padding token predictions.

Because we don't care about further prediction after the pad_token is predicted for the first time, so it doesn't make sense to punish our network after the pad_token is predicted.

For simplicity you can ignore this comment, it's up to you.

```
In [17]: 1 # Define the loss as categorical cross-entropy (e.g. from keras.losses).
2 # Mind that predictions are probabilities and NOT logits!
3 # Remember to apply tf.reduce_mean to get a scalar loss!
4 ### YOUR CODE HERE
5 loss = tf.reduce_mean(tf.keras.losses.categorical_crossentropy(answers_matrix, predictions_matrix))
6
7 optimize = tf.train.AdamOptimizer().minimize(loss)
```


WARNING:tensorflow:From C:\Users\Xiaowei\Anaconda3\envs\tfspark\lib\site-packages\tensorflow\python\ops\math_grad.py:1250: add_dispatch_support.<locals>.wrapper (from tensorflow.python.ops.array_ops) is deprecated and will be removed in a future version.

Instructions for updating:

Use tf.where in 2.0, which has the same broadcast rule as np.where

RNN: training

```
In [18]: 1 from IPython.display import clear_output
2 from random import sample
3
4 s.run(tf.global_variables_initializer())
5
6 batch_size = 32
7 history = []
8
9 for i in range(1000):
10     batch = to_matrix(sample(names, batch_size), max_len=MAX_LENGTH)
11     loss_i, _ = s.run([loss, optimize], {input_sequence: batch})
12
13     history.append(loss_i)
14
15     if (i + 1) % 100 == 0:
16         clear_output(True)
17         plt.plot(history, label='loss')
18         plt.legend()
19         plt.show()
20
21 assert np.mean(history[:10]) > np.mean(history[-10:]), "RNN didn't converge"
```



RNN: sampling

Once we've trained our network a bit, let's get to actually generating stuff. All we need is the `rnn_one_step` function you have written above.

```
In [19]: 1 x_t = tf.placeholder(tf.int32, (1,))
2 h_t = tf.Variable(np.zeros([1, rnn_num_units], np.float32)) # we will update hidden state in this variable
3
4 # For sampling we need to define `rnn_one_step` tensors only once in our graph.
5 # We reuse all parameters thanks to functional API usage.
6 # Then we can feed appropriate tensor values using feed_dict in a loop.
7 # Note how different it is from training stage, where we had to unroll the whole sequence for backprop.
8 next_probs, next_h = rnn_one_step(x_t, h_t)
```

```
In [20]: 1 def generate_sample(seed_phrase=start_token, max_length=MAX_LENGTH):
2     '''
3     This function generates text given a `seed_phrase` as a seed.
4     Remember to include start_token in seed phrase!
5     Parameter `max_length` is used to set the number of characters in prediction.
6     '''
7     x_sequence = [token_to_id[token] for token in seed_phrase]
8     s.run(tf.assign(h_t, h_t.initial_value))
9
10    # feed the seed phrase, if any
11    for ix in x_sequence[:-1]:
12        s.run(tf.assign(h_t, next_h), {x_t: [ix]})
13
14    # start generating
15    for _ in range(max_length-len(seed_phrase)):
16        x_probs, _ = s.run([next_probs, tf.assign(h_t, next_h)], {x_t: [x_sequence[-1]]})
17        x_sequence.append(np.random.choice(n_tokens, p=x_probs[0]))
18
19    return ''.join([tokens[ix] for ix in x_sequence if tokens[ix] != pad_token])
```

In [21]:

```
1 # without prefix
2 for _ in range(10):
3     print(generate_sample())
```

Lhanla
Delse
Aleree
Amrin
Asdiy
Larrea
Cylxep
Sribertir
Andie
Kabyr

In [22]:

```
1 # with prefix conditioning
2 for _ in range(10):
3     print(generate_sample(' Trump'))
```

Trumpadno
Trumpa
Trumpis
Trumpie
Trumpa
Trumpenda
Trumpie
Trumpand
Trumpada
Trumpar

Submit to Coursera

In [23]:

```
1 # token expires every 30 min
2 COURSERA_TOKEN = "e2f2dGXhTPjAMzyJ"
3 COURSERA_EMAIL = "lxwvictor@gmail.com"
```

In [24]:

```
1 from submit import submit_char_rnn
2 samples = [generate_sample(' AI') for i in tqdm_utils.tqdm_notebook_failsafe(range(25))]
3 submission = (history, samples)
4 submit_char_rnn(submission, COURSERA_EMAIL, COURSERA_TOKEN)
```

100%

25/25 [00:19<00:00, 1.29it/s]

You used an invalid email or your token may have expired. Please make sure you have entered all fields correctly. Try generating a new token if the issue still persists.

Try it out!

Disclaimer: This part of assignment is entirely optional. You won't receive bonus points for it. However, it's a fun thing to do. Please share your results on course forums.

You've just implemented a recurrent language model that can be tasked with generating any kind of sequence, so there's plenty of data you can try it on:

- Novels/poems/songs of your favorite author
- News titles/clickbait titles
- Source code of Linux or Tensorflow
- Molecules in [smiles \(https://en.wikipedia.org/wiki/Simplified_molecular-input_line-entry_system\)](https://en.wikipedia.org/wiki/Simplified_molecular-input_line-entry_system) format
- Melody in notes/chords format
- IKEA catalog titles
- Pokemon names
- Cards from Magic, the Gathering / Hearthstone

If you're willing to give it a try, here's what you wanna look at:

- Current data format is a sequence of lines, so a novel can be formatted as a list of sentences. Alternatively, you can change data preprocessing altogether.
- While some datasets are readily available, others can only be scraped from the web. Try Selenium or Scrapy for that.
- Make sure MAX_LENGTH is adjusted for longer datasets. There's also a bonus section about dynamic RNNs at the bottom.
- More complex tasks require larger RNN architecture, try more neurons or several layers. It would also require more training iterations.
- Long-term dependencies in music, novels or molecules are better handled with LSTM or GRU

Good hunting!

Bonus level: dynamic RNNs

Apart from Keras, there's also a friendly TensorFlow API for recurrent neural nets. It's based around the symbolic loop function (aka [tf.scan](https://www.tensorflow.org/api_docs/python/tf/scan) (https://www.tensorflow.org/api_docs/python/tf/scan)).

RNN loop that we implemented for training can be replaced with single TensorFlow instruction: [tf.nn.dynamic_rnn](https://www.tensorflow.org/api_docs/python/tf/nn/dynamic_rnn) (https://www.tensorflow.org/api_docs/python/tf/nn/dynamic_rnn). This interface allows for dynamic sequence length and comes with some pre-implemented architectures.

Take a look at [tf.nn.rnn_cell.BasicRNNCell](https://www.tensorflow.org/api_docs/python/tf/contrib/rnn/BasicRNNCell) (https://www.tensorflow.org/api_docs/python/tf/contrib/rnn/BasicRNNCell).

```
In [ ]: 1 class CustomRNN(tf.nn.rnn_cell.BasicRNNCell):
2         def call(self, input, state):
3             # from docs:
4             # Returns:
5             # Output: A 2-D tensor with shape [batch_size, self.output_size].
6             # New state: Either a single 2-D tensor, or a tuple of tensors matching the arity and shapes of state.
7             return rnn_one_step(input[:, 0], state)
8
9         @property
10        def output_size(self):
11            return n_tokens
12
13    cell = CustomRNN(rnn_num_units)
14
15    input_sequence = tf.placeholder(tf.float32, (None, None))
16
17    predicted_probas, last_state = tf.nn.dynamic_rnn(cell, input_sequence[:, :, None], dtype=tf.float32)
18
19    sess = tf.Session()
20    init = tf.global_variables_initializer()
21    sess.run(init)
22    with sess.as_default():
23        print('LSTM outputs for each step [batch,time,n_tokens]:')
24        print(predicted_probas.eval({input_sequence: to_matrix(names[:10], max_len=50)}).shape)
```

Note that we never used MAX_LENGTH in the code above: TF will iterate over however many time-steps you gave it.

You can also use any pre-implemented RNN cell:

```
In [ ]: 1 for obj in dir(tf.nn.rnn_cell) + dir(tf.contrib.rnn):
2         if obj.endswith('Cell'):
3             print(obj, end="\t")
```

```
In [ ]: 1 input_sequence = tf.placeholder(tf.int32, (None, None))
2
3     inputs_embedded = embed_x(input_sequence)
4
5     # standard cell returns hidden state as output!
6     cell = tf.nn.rnn_cell.LSTMCell(rnn_num_units)
7
8     state_sequence, last_state = tf.nn.dynamic_rnn(cell, inputs_embedded, dtype=tf.float32)
9
10    s.run(tf.global_variables_initializer())
11
12    print('LSTM hidden state for each step [batch,time,rnn_num_units]:')
13    print(state_sequence.eval({input_sequence: to_matrix(names[:10], max_len=50)}).shape)
```

```
In [ ]: 1
```